

random-forest

JackMoorer

```
train <- read.csv("/Users/jackmoorer/Stat154/Projects/Project/data/clean_train.csv", header = TRUE)
test <- read.csv("/Users/jackmoorer/Stat154/Projects/Project/data/clean_test.csv", header = TRUE)
```

```
library(ggplot2)
library(randomForest)
```

```
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
library(caret)
```

```
## Loading required package: lattice
```

Train Set

```
train_forest <- train[, -ncol(train)]
```

First I tried just the basic random forest to see the results

```
set.seed(200)
rf <- randomForest(Over50k ~., data = train_forest, importance = TRUE)
```

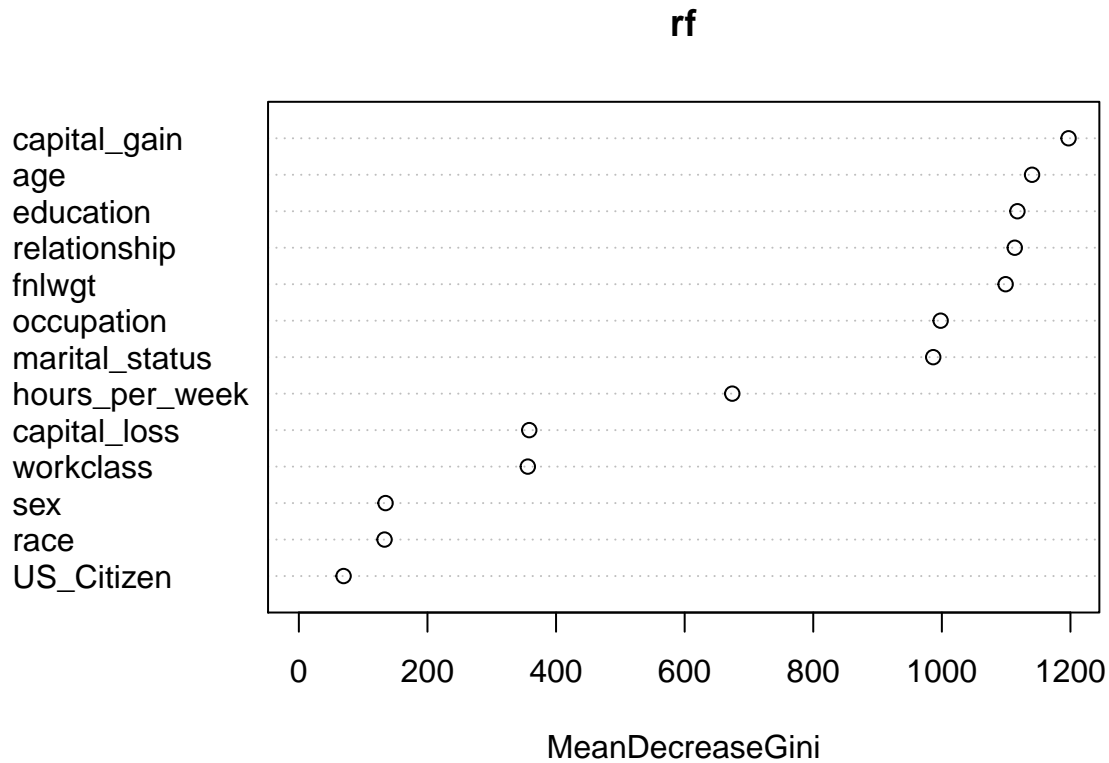
```
rf$confusion
```

```
##           No  Yes class.error
## No   21177 1477   0.0651982
## Yes   2622 4886   0.3492275
```

```
importance(rf)
```

```
##           MeanDecreaseGini
## age                1140.26647
## workclass           356.09448
## fnlwt              1099.15031
## education           1117.53836
## marital_status       986.55330
## occupation           998.11101
## relationship         1113.39050
## race                 133.24092
## sex                  134.81105
## capital_gain         1197.05878
## capital_loss          358.29773
## hours_per_week        674.04817
## US_Citizen            69.55423
```

```
varImpPlot(rf)
```



```
train_predictors <- train_forest[, -ncol(train_forest)]
```

```
train_response <- train_forest$Over50k
```

Next I did cross validation to tune the number of predictors used while building a tree


```
set.seed(200)
```

```
rf_cv <- rfcv(train_predictors, train_response, cv.fold = 5)
```

The plot below shows the best number of predictors to use is all 13.

```
with(rf_cv, plot(n.var, error.cv, log="x", type="o", lwd=1, main = "CV Variables vs Error Rate", xlab =
```

CV Variables vs Error Rate



```
rf_cv$error.cv[[1]]
```

```
## [1] 0.1367615
```

```
num_var <- as.numeric(names(which.min(rf_cv$error.cv)))
print(num_var)
```

```
## [1] 13
```

```
#this takes too long
```

```
set.seed(200)
ntree <- c(50, 100, 500, 1000, 1500)
matrix <- matrix(rep(0, 4*5*3), ncol = 3, nrow = 20)
i <- 0
for (tree in ntree) {
  rf_cv <- rfcv(train_predictors, train_response, cv.fold = 3, ntree = tree)
  for (j in 1:length(rf_cv$n.var)) {
    matrix[i + j, ] <- c(tree, rf_cv$n.var[j], rf_cv$error.cv[[j]])
  }
  i <- i + j
}
```

```
cv_df <- data.frame(matrix)
names(cv_df) <- c("ntree", "num_var", "error_rate")
err_rate <- cv_df$error_rate
row_num <- as.numeric(which.min(err_rate))
best_ntree <- cv_df$ntree[row_num]
best_num_var <- cv_df$num_var[row_num]
```

```
set.seed(200)
ntree <- c(50, 100, 500, 1000, 1500)
cv_list <- as.list(rep(0, 5))
```

```

names(cv_list) <- ntree
i <- 1
for (tree in ntree) {
  tune_param <- tuneRF(train_predictors, train_response, ntreeTry = tree, trace = FALSE, plot = FALSE)
  cv_list[[i]] <- tune_param
  i <- i + 1
}

```

```

## 0.003301108 0.05
## -0.06366423 0.05
## -0.011597 0.05
## -0.05146171 0.05
## -0.0004873294 0.05
## -0.05165692 0.05
## 0.008464329 0.05
## -0.04353083 0.05
## 0.002679006 0.05
## -0.04627375 0.05

```

```
cv_list
```

```

## $`50`
##      mtry  OOBError
## 2.00B    2 0.1401432
## 3.00B    3 0.1406074
## 6.00B    6 0.1495590
##
## $`100`
##      mtry  OOBError
## 2.00B    2 0.1388171
## 3.00B    3 0.1372256
## 6.00B    6 0.1442875
##
## $`500`
##      mtry  OOBError
## 2.00B    2 0.1361316
## 3.00B    3 0.1360652
## 6.00B    6 0.1430940
##
## $`1000`
##      mtry  OOBError
## 2.00B    2 0.1359326
## 3.00B    3 0.1370930
## 6.00B    6 0.1430608
##
## $`1500`
##      mtry  OOBError
## 2.00B    2 0.1357669
## 3.00B    3 0.1361316
## 6.00B    6 0.1424309

```

```

library(stringr)
matrix <- matrix(rep(0, 3*5*3), nrow = 15, ncol = 3)
index <- 0
for (i in 1:5) {

```

```

cur_cv <- cv_list[i]
errs <- cv_list[[i]][,2]
for (j in 1:length(errs)){
  cur_err <- cv_list[[i]][,2][[j]]
  val <- str_sub(names(cv_list[[i]][,2][j]), 1, 1)
  val <- as.numeric(val)
  matrix[index + j, ] <- c(ntree[i], val, cur_err)
}
index <- index + j
}
cv_df <- data.frame(matrix)
names(cv_df) <- c("ntree", "num_var", "OOB")
cv_df

```

```

##      ntree num_var      OOB
## 1      50      2 0.1401432
## 2      50      3 0.1406074
## 3      50      6 0.1495590
## 4     100      2 0.1388171
## 5     100      3 0.1372256
## 6     100      6 0.1442875
## 7     500      2 0.1361316
## 8     500      3 0.1360652
## 9     500      6 0.1430940
## 10    1000      2 0.1359326
## 11    1000      3 0.1370930
## 12    1000      6 0.1430608
## 13    1500      2 0.1357669
## 14    1500      3 0.1361316
## 15    1500      6 0.1424309

```

```

library(ggplot2)
cv_df$number_of_trees <- factor(cv_df$ntree)


```

```

ggplot(data = cv_df, aes(x = num_var, y = OOB, color = number_of_trees)) + geom_line() + ggtitle("Random

```

Random Forest Hyper Parameter Tuning



```
oobs <- cv_df$OOB
row_num <- as.numeric(which.min(oobs))
num_tree <- cv_df[row_num, 1]
num_var <- cv_df[row_num, 2]
```

I am going to run the random forest using this parameter:

```
set.seed(100)
random_forest <- randomForest(Over50k ~ ., data = train_forest, ntree = num_tree, mtry = num_var, importance = TRUE)
print(random_forest)
```

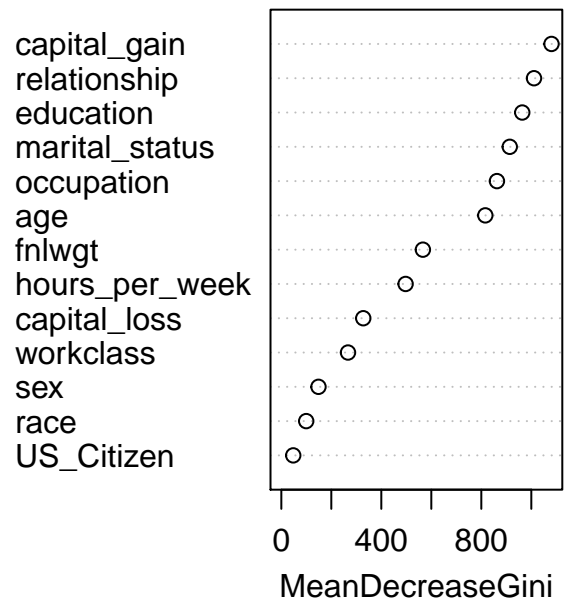
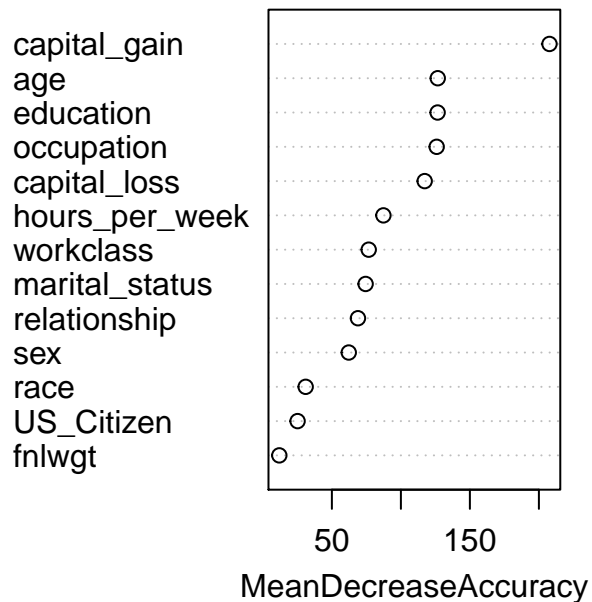
```
##
## Call:
## randomForest(formula = Over50k ~ ., data = train_forest, ntree = num_tree, mtry = num_var, importance = TRUE)
##              Type of random forest: classification
##              Number of trees: 1500
## No. of variables tried at each split: 2
##
## OOB estimate of  error rate: 13.63%
## Confusion matrix:
##      No  Yes class.error
## No 21308 1346 0.05941556
## Yes 2764 4744 0.36814065
importance(random_forest)
```

```
##              No              Yes MeanDecreaseAccuracy
## age          2.916181 117.7812947          126.71208
```

## workclass	67.458005	27.2146935	76.69084
## fnlwgt	13.726050	0.5755618	11.91015
## education	82.102928	96.6243220	126.61319
## marital_status	75.476360	44.6633592	74.44127
## occupation	70.493598	96.1170909	125.93797
## relationship	45.553450	66.4228965	68.89658
## race	19.108144	20.2750338	31.00323
## sex	44.129040	10.9439606	62.22469
## capital_gain	172.812802	216.5955580	207.66669
## capital_loss	88.886756	115.6337429	117.23932
## hours_per_week	18.625739	84.8279608	87.33325
## US_Citizen	27.569738	1.2870478	25.15558
##	MeanDecreaseGini		
## age		815.99695	
## workclass		266.76014	
## fnlwgt		566.25831	
## education		963.76462	
## marital_status		913.81883	
## occupation		862.54477	
## relationship		1010.80850	
## race		99.57915	
## sex		148.65378	
## capital_gain		1080.78006	
## capital_loss		327.79417	
## hours_per_week		497.17865	
## US_Citizen		47.64475	

```
varImpPlot(random_forest)
```

random_forest



```
rf_pred <- predict(random_forest, train_forest[, -ncol(train_forest)])
```

```
real_train <- train_forest$Over50k  
err_rate <- mean(rf_pred != real_train)  
err_rate
```

```
## [1] 0.08938399
```

```
library(ROCR)
```

```
## Loading required package: gplots
```

```
##
```

```
## Attaching package: 'gplots'
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
## lowess
```

```
probs <- as.vector(random_forest$votes[,2])
```


```
prediction <- prediction(probs, real_train)
```

```
performance <- performance(prediction, measure = "tpr", x.measure = "fpr")
```

```
plot(performance, main="ROC Random Forest")
```

```
abline(a=0, b=1, lty = 2)
```

ROC Random Forest



```
auc <- performance(prediction, measure="auc")@y.values[[1]]
```

```
auc
```

```
## [1] 0.9071512
```



```
print(random_forest$confusion)
```

```
##           No  Yes class.error
## No  21308 1346  0.05941556
## Yes   2764 4744  0.36814065
```

Test Set

```
test_forest <- test[, -ncol(test)]
```

```
test_forest_preds <- test_forest[, -ncol(test_forest)]
real_test <- test_forest$Over50k
```

```
rf_test_preds <- predict(random_forest, test_forest_preds)
err_rate <- mean(rf_test_preds != real_test)
err_rate
```

```
## [1] 0.1394422
```

```
confusionMatrix <- table(rf_test_preds, real_test)
confusionMatrix
```

```
##           real_test
## rf_test_preds    No  Yes
##           No 10638 1378
##           Yes   722 2322
```

```
sensitivity <- confusionMatrix[2, 2]/(confusionMatrix[2, 2] + confusionMatrix[1, 2])
sensitivity
```

```
## [1] 0.6275676
```

```
specificity <- confusionMatrix[1, 1]/(confusionMatrix[1, 1] + confusionMatrix[2, 1])
specificity
```


```
## [1] 0.9364437
```

```
test_probs <- predict(random_forest, test_forest_preds, type = "prob")
```

```
test_prediction <- prediction(test_probs[,2], real_test)
test_performance <- performance(test_prediction, measure = "tpr", x.measure = "fpr")
```

```
plot(test_performance, main="Test ROC Random Forest")
abline(a=0, b=1, lty=2)
```

Test ROC Random Forest



```
auc <- performance(test_prediction, measure="auc")@y.values[[1]]
auc
```

```
## [1] 0.9031879
```